
rl-triton: High-Performance Triton GPU Kernels for Reinforcement Learning Credit Assignment

A PREPRINT

Lars Simon Zehnder*
Independent Researcher
Faro, Portugal

ABSTRACT

We present `rl-triton`, an open-source library of high-performance GPU kernels for reinforcement learning credit assignment, implemented in Triton. The core contribution is a unified associative scan framework that recasts seven distinct RL estimation algorithms – Generalized Advantage Estimation (GAE), V-Trace, Retrace(λ), TD(λ) returns, discounted returns, eligibility traces, and episodic prefix sums – as instances of a single first-order linear recurrence solved in $O(\log T)$ parallel steps. All algorithms share the same associative scan operator, with algorithm-specific fused Triton kernels constructing their recurrence coefficients on-chip. We verify the associative operator algebraically and define the treatment of terminated and truncated episodes explicitly. Benchmarks show a $1.6\text{--}5.70\times$ full-call speedup over a vectorized `torch.compile` baseline in the massively parallel simulation regime (thousands of environments, short rollouts). The reported range covers all seven algorithms on both GPUs, both with and without per-step truncation handling. For most algorithms, speedups increase at longer sequence lengths, where the baseline incurs more $O(\log T)$ HBM round-trips. The library is available at <https://github.com/simonsays1980/rl-triton>.

1 Introduction

Every iteration of a deep reinforcement learning training pipeline requires computing *credit assignment* quantities: advantages, value targets, and return estimates that tell the policy network which actions were good and by how much. These quantities include Generalized Advantage Estimation (GAE) [Schulman et al., 2016], V-Trace targets [Espeholt et al., 2018], Retrace(λ) Q-value targets [Munos et al., 2016], and several simpler estimators. While network forward and backward passes and environment rollouts typically dominate overall training time, credit assignment computations are on the critical path of every update step. They operate over tensors of shape `[num_envs, seq_len]`, where `num_envs` is the number of parallel environments (or workers) and `seq_len=T` is the number of timesteps in each rollout-buffer row, i.e., the scan length. It is not necessarily the episode length, since a row may contain multiple episodes separated by boundary flags. Typical configurations range from a handful of environments with short rollouts (e.g., 8 environments, 128 steps in standard PPO [Schulman et al., 2017]) to thousands of environments with longer horizons in large-scale distributed training [Espeholt et al., 2018, Petrenko et al., 2020] and RLHF pipelines where rollout sequences span thousands of tokens [Ouyang et al., 2022].

All of these estimators share a common mathematical structure: they are **first-order linear recurrences**. The five backward estimators have the form

$$A_t = \alpha_t + \beta_t \cdot A_{t+1}, \quad A_T = 0 \text{ (or a bootstrap value)}, \quad (1)$$

where α_t and β_t are algorithm-specific per-step coefficients. Eligibility traces and episodic prefix sums use the forward-time mirror $A_t = \alpha_t + \beta_t A_{t-1}$. In either direction, the recurrence creates a strict temporal dependency chain: a naive sequential implementation requires T serial steps regardless of how many parallel GPU cores are available.

*Correspondence: `simon@neway.ai`, `simon.zehnder@gmail.com`. Code: <https://github.com/simonsays1980/rl-triton>

Existing RL codebases commonly evaluate this recurrence with a sequential loop over t (Section 8). Wrapping that loop in `torch.compile` removes Python overhead but does not remove its temporal dependency. An $O(\log T)$ restructuring instead requires expressing the recurrence as an explicit parallel associative scan [Blelloch, 1990] (Section 2.2). We implement this construction in PyTorch as the baseline used in our evaluation (Section 6.2). Although it has the same asymptotic depth as the Triton scan, each doubling step still incurs an HBM round-trip.

We implement fused Triton kernels for these recurrences. The five backward algorithms use Equation 1, while eligibility traces and episodic prefix sums use its forward-time mirror. Each algorithm has its own fused kernel, which constructs its coefficients (Section 3) from raw rollout inputs on-chip and evaluates the recurrence using the same associative scan operator and combine function.

Contributions.

1. A unified algebraic framework showing that seven standard RL credit assignment algorithms are instances of a single first-order linear recurrence.
2. Fused Triton kernels sharing one associative scan, with on-chip coefficient construction and explicit termination/truncation handling.
3. Algorithm-specific treatment of bootstrap values at rollout-window boundaries – an additive term in α_{T-1} for some algorithms and a nonzero scan carry for others (Section 3.2).
4. A benchmarking harness showing a 1.6–5.70 $\times$ full-call speedup over a vectorized `torch.compile` baseline across all seven algorithms and both GPUs in the massively parallel simulation regime (Section 7).
5. An open-source Python package with correctness tests, performance regression tests, and documentation.

2 Background

2.1 The Sequential Dependency Problem

Consider the GAE recurrence [Schulman et al., 2016]:

$$A_t = \delta_t + \gamma \lambda \cdot A_{t+1},$$

where $\delta_t = r_t + \gamma V(s_{t+1}) - V(s_t)$ is the one-step TD error. Computing A_0 requires A_1 , which requires A_2 , all the way to A_{T-1} . This is a strict sequential chain of length T .

On GPU hardware, sequential chains of length T are expensive not because of arithmetic cost but because of **memory traffic**: a naive loop must read δ_t and $\gamma \lambda$ from HBM, perform the addition and multiplication, and write A_t back to HBM – one full round-trip per timestep. For $T=1024$, `num_envs=128`, this means 1024 serial HBM round-trips even though 128 environment rows could in principle be processed in parallel.

`torch.compile` cannot remove this dependency: it eliminates per-step Python interpreter overhead, but a compiler cannot break a true data dependency the algorithm itself imposes, so the loop still executes T sequential steps. This uncompiled sequential loop – the pattern used by CleanRL, RLlib, and similar codebases (Section 8) – is reported as the *Loop* baseline in our results (Section 7).

Removing the $O(T)$ chain requires restructuring the *algorithm*, not just compiling it: rewriting the recurrence as an associative scan (Section 2.2 below) so that `torch.compile` has a $\log_2 T$ -depth computation to compile in the first place, rather than a T -depth one. This is the vectorized, `torch.compile`’d doubling-scan baseline used throughout this paper’s evaluation (Section 6.2) and is a substantially stronger baseline than the naive loop – it achieves the same $O(\log T)$ depth as our Triton kernel’s internal scan (Section 4.2) – but, as Section 6.2 details, still pays a full HBM round-trip per doubling step rather than keeping intermediate results on-chip.

2.2 Parallel Associative Scans

The parallel associative scan is a foundational algorithm in parallel computing [Blelloch, 1990, Ladner and Fischer, 1980]. Given an array of elements $[x_0, x_1, \dots, x_{T-1}]$ and an associative binary operator $\oplus$, the scan computes all prefix reductions $[x_0, x_0 \oplus x_1, \dots, x_0 \oplus \dots \oplus x_{T-1}]$ in $O(\log T)$ parallel steps on T processors, using a binary tree reduction.

The first-order linear recurrence $A_t = \alpha_t + \beta_t \cdot A_{t+1}$ can be expressed as a prefix scan of tuples (α_t, β_t) under the associative operator:

$$(\alpha_B, \beta_B) \oplus (\alpha_A, \beta_A) = (\alpha_B + \beta_B \alpha_A, \beta_A \beta_B). \quad (2)$$

Proposition 1 (Associativity of $\oplus$). *The operator defined by Equation 2 is associative.*

Proof. For any three tuples $(\alpha_C, \beta_C), (\alpha_B, \beta_B), (\alpha_A, \beta_A)$:

$$\begin{aligned} [(\alpha_C, \beta_C) \oplus (\alpha_B, \beta_B)] \oplus (\alpha_A, \beta_A) &= (\alpha_C + \beta_C \alpha_B, \beta_B \beta_C) \oplus (\alpha_A, \beta_A) \\ &= (\alpha_C + \beta_C \alpha_B + \beta_B \beta_C \alpha_A, \beta_A \beta_B \beta_C). \\ (\alpha_C, \beta_C) \oplus [(\alpha_B, \beta_B) \oplus (\alpha_A, \beta_A)] &= (\alpha_C, \beta_C) \oplus (\alpha_B + \beta_B \alpha_A, \beta_A \beta_B) \\ &= (\alpha_C + \beta_C (\alpha_B + \beta_B \alpha_A), \beta_A \beta_B \beta_C) \\ &= (\alpha_C + \beta_C \alpha_B + \beta_C \beta_B \alpha_A, \beta_A \beta_B \beta_C). \end{aligned}$$

Both expressions are equal. $\square$

For the five backward recurrences, the kernel loads the sequence in reverse chronological order: reversed index $k = 0$ corresponds to timestep $t = T - 1$, while $k = T - 1$ corresponds to $t = 0$. The associative scan therefore operates in late-to-early time order, and the result at reversed index k corresponds to chronological timestep $t = T - 1 - k$. Eligibility traces and episodic prefix sums instead scan the sequence in its original chronological order.

2.3 Triton and the GPU Memory Hierarchy

Triton [Tillet et al., 2019] is a domain-specific language and compiler for writing GPU kernels in Python-like syntax. It abstracts over warp-level programming while exposing the key performance-relevant features of the GPU memory hierarchy:

- **HBM (High Bandwidth Memory):** Main GPU memory, $\sim 2\text{--}3.35$ TB/s bandwidth on datacenter HBM3 parts such as the 80GB H100 used in this paper’s evaluation, latency on the order of several hundred cycles. Accessible by all SMs (The RTX 2000 Ada also evaluated here uses 16GB of GDDR6 instead of HBM, at substantially lower bandwidth – Section 7 reports results on both GPUs).
- **SRAM (Shared Memory):** Per-SM scratchpad, $\sim 5\text{--}10\times$ higher bandwidth than HBM, very low latency. Configurable capacity up to ~ 100 KB per SM on Ada-generation SMs (RTX 2000 Ada) and up to ~ 228 KB per SM on H100.
- **Registers:** Per-thread storage, zero-latency access. Arithmetic occurs here. 64K 32-bit registers per SM on both GPUs evaluated in this paper, shared across all resident threads.

Credit assignment kernels are memory-bandwidth bound rather than compute bound [Williams et al., 2009]: the arithmetic intensity (FLOPs per byte of HBM traffic) is low, so reducing HBM accesses directly reduces runtime. The performance advantage of a fused kernel therefore comes from **minimizing HBM round-trips** [Dao et al., 2022]. A fused implementation keeps the scan state on-chip between stages, avoiding the intermediate HBM round-trips required by a materialized scan. The associative scan keeps intermediate tuples on-chip throughout the $O(\log T)$ reduction, with no HBM round-trip until the final output store.

Triton’s `tl.associative_scan` primitive maps directly to this pattern, allowing the kernel to express the tree reduction at a high level while the Triton compiler generates efficient hardware-specific reduction instructions for both NVIDIA (warp shuffles) and AMD (wavefront operations) GPUs.

3 The Unified Linear Recurrence Framework

3.1 The Shared Scan Operator

All seven algorithms use the same associative scan operator and combine function. For the five backward recurrences, the kernel presents (α_t, β_t) to the scan in reverse chronological order and writes the result back in original time order. Eligibility traces and episodic prefix sums instead scan the original sequence in forward order.

For the backward algorithms, the boundary carry A_T is algorithm-dependent rather than a universal constant; Section 3.2 derives this condition, and Table 1 lists it for each algorithm. The scan incorporates this carry after evaluating the recurrence with `tl.associative_scan` and the combine function in Equation 2. The forward scans use the same combine function but require no successor-value bootstrap.

Table 1: Algorithm-specific coefficients for the shared affine scan, applied in reverse or forward time as appropriate. $d_t^{\text{term}}, d_t^{\text{trunc}} \in \{0, 1\}$ flag termination and truncation; $d_t = d_t^{\text{term}} \vee d_t^{\text{trunc}}$ is the flag used in β_t , combining both. $\Delta_t = v_t - V(s_t)$ is the V-Trace value delta (Section 3.3.2). The **Bootstrap** column gives each algorithm’s scan boundary carry (Section 3.2 explains why it is 0 for GAE/V-Trace/Retrace(λ) but nonzero for TD(λ)/discounted returns). Eligibility traces use vector notation since the trace is vector-valued; see Section 9.2 for per-call scalar handling.

Algorithm	α_t	β_t	Bootstrap
GAE	$r_t + \gamma(1 - d_t^{\text{term}})V(s_{t+1}) - V(s_t)$	$\gamma\lambda(1 - d_t)$	$A_T = 0$
V-Trace	$\rho_t(r_t + \gamma(1 - d_t^{\text{term}})V(s_{t+1}) - V(s_t))$	$\gamma c_t(1 - d_t)$	$\Delta_T = 0$
Retrace(λ)	$r_t + \gamma(1 - d_t^{\text{term}})\mathbb{E}_\pi[Q(s_{t+1}, \cdot)] - Q(s_t, a_t)$	$\gamma c_{t+1}(1 - d_t)$	0
TD(λ)	$r_t + \gamma(1 - \lambda)(1 - d_t^{\text{term}})V(s_{t+1})$	$\gamma\lambda(1 - d_t)$	G_T^λ
Disc. Returns	r_t	$\gamma(1 - d_t)$	G_T
Elig. Traces (fwd)	$\nabla_{\mathbf{w}} \hat{V}(s_t, \mathbf{w}_t)$	$\gamma\lambda(1 - d_{t-1})$	$\mathbf{z}_{-1} = \mathbf{0}$
Prefix Sum (fwd)	x_t	$1 - d_{t-1}$	0

For the common case (`seq_len` $\leq$ 131072), each Python API dispatches to an algorithm-specific fused kernel that receives the raw rollout tensors (rewards, values, done flags, importance ratios, etc.) and constructs α_t and β_t in registers; no PyTorch operation materializes them as intermediate tensors. For longer sequences, the backward-scan algorithms use the unfused chunked fallback described in Section 4.3. Table 1 summarizes the per-algorithm mapping from rollout quantities to α_t and β_t .

3.2 Episode and Rollout-Window Boundary Semantics

Credit-assignment quantities must not propagate across episode boundaries within a rollout buffer. `rl-triton` distinguishes termination and truncation, which both sever the scan carry but differ in their bootstrap semantics, from the rollout-window boundary at $t = T - 1$. Eligibility traces and episodic prefix sums do not use value bootstrapping; for these algorithms, episode boundaries only reset the recurrence.

Terminated steps ($d_t^{\text{term}} = 1$). The episode ended naturally. For algorithms that bootstrap from a successor value, $V(s_{t+1})$ is conceptually zero and is therefore removed from the local update through $(1 - d_t^{\text{term}})$. The decay β_t is also zeroed through $(1 - d_t)$, preventing the scan carry from crossing into the next episode.

Truncated steps ($d_t^{\text{trunc}} = 1$). The episode was cut short but continues in principle. The value stored at $t + 1$ belongs to the next episode in the rollout buffer and therefore cannot serve as the continuation value. For algorithms that bootstrap from $V(s_{t+1})$, the caller instead supplies the correct continuation through `bootstrap_values[env, t]`; this contribution is incorporated into α_t with the appropriate discount. As for termination, $\beta_t = 0$ severs the scan carry, but does not remove the bootstrap already encoded in α_t .

Rollout-window boundary (final step $t = T - 1$). If the rollout ends in the middle of an episode, its successor state s_T lies outside the buffer. For GAE, V-Trace, and Retrace(λ), the supplied $V(s_T)$ enters the local update at $t = T - 1$, while the scan boundary carry remains

$$A_T = 0.$$

A nonzero carry would represent continuation of the trace recurrence beyond the sampled window; the required one-step bootstrap has already been included in α_{T-1} . For TD(λ) and discounted returns, the supplied continuation value at s_T is instead distributed between α_{T-1} and the boundary carry A_T , with coefficients that sum to the required discounted continuation (Section 3.3.4). Eligibility traces and episodic prefix sums do not require a successor-value bootstrap.

If $t = T - 1$ is itself truncated, it is handled as a truncation rather than as an ordinary non-terminal window edge: the supplied continuation value enters through the truncation term in α_{T-1} , while $\beta_{T-1} = 0$ suppresses any contribution from the scan carry. The bootstrap is therefore not double-counted.

Thus, termination and truncation both reset the recurrence, but termination uses a zero successor value whereas truncation uses a supplied continuation value. A non-terminal rollout-window edge does not represent an episode boundary; its continuation value enters through α_{T-1} , the boundary carry A_T , or both, depending on the algorithm (Section 3.3).

3.3 Algorithm-Specific Details

The following subsections specify α_t and β_t for each algorithm. Episode-boundary semantics – how terminated steps, truncated steps, and the window boundary are handled – are common to all algorithms and described in Section 3.2; they are not repeated here.

3.3.1 Generalized Advantage Estimation (GAE)

GAE [Schulman et al., 2016] computes the advantage as the exponentially-weighted mixture of n -step TD residuals:

$$A_t^{\text{GAE}} = \sum_{k=0}^{\infty} (\gamma\lambda)^k \delta_{t+k}.$$

This telescopes to the backward recurrence $A_t = \delta_t + \gamma\lambda \cdot A_{t+1}$. With episode-boundary masking:

$$\alpha_t = \delta_t = r_t + \gamma(1 - d_t^{\text{term}})V(s_{t+1}) - V(s_t), \quad \beta_t = \gamma\lambda(1 - d_t).$$

The $\lambda = 0$ case reduces GAE to a pure one-step TD advantage (low variance, high bias); $\lambda = 1$ yields the full Monte Carlo return minus the baseline (low bias, high variance).

3.3.2 V-Trace

V-Trace [Espeholt et al., 2018] corrects for policy lag in distributed training by weighting each TD error with a clipped importance sampling (IS) ratio. Let π be the current learner policy and μ the behaviour policy that collected the data. Define the per-step IS ratio $\rho_t^{\text{raw}} = \pi(a_t|s_t)/\mu(a_t|s_t)$, then clip it at two separate thresholds:

$$\rho_t = \min(\bar{\rho}, \rho_t^{\text{raw}}), \quad c_t = \min(\bar{c}, \rho_t^{\text{raw}}).$$

$\bar{\rho}$ controls the bias–variance trade-off of the value target: a smaller $\bar{\rho}$ limits how far the target can deviate from the on-policy value, reducing variance at the cost of some bias. $\bar{c}$ controls the multi-step trace length: it multiplies along the return trajectory in $\beta_t = \gamma c_t(1 - d_t)$, so a smaller $\bar{c}$ shortens the effective lookahead, trading variance for bias. Typical defaults are $\bar{\rho} = 1$ and $\bar{c} = 1$ [Espeholt et al., 2018]. The V-Trace target delta satisfies $\Delta_t = \delta_t^V + \gamma c_t \cdot \Delta_{t+1}$, mapping to:

$$\alpha_t = \delta_t^V = \rho_t(r_t + \gamma(1 - d_t^{\text{term}})V(s_{t+1}) - V(s_t)), \quad \beta_t = \gamma c_t(1 - d_t).$$

After the scan, targets are formed as $v_t = \Delta_t + V(s_t)$ and advantages as $A_t = \rho_t(r_t + \gamma v_{t+1} - V(s_t))$, all computed within the same fused kernel.

3.3.3 Retrace(λ)

Retrace(λ) [Munos et al., 2016] extends V-Trace to Q-value estimation with an *index shift*: the decay coefficient applying to Δ_{t+1} is c_{t+1} (next-step importance weight) rather than c_t :

$$\Delta_t = \delta_t + \gamma c_{t+1} \cdot \Delta_{t+1}.$$

This shift reflects that Retrace evaluates a state-action pair (s_t, a_t) : the action a_t is fixed as the subject of evaluation, so off-policy correction begins only from the next action a_{t+1} onward. Because the kernel already loads (α_t, β_t) in reverse chronological order (Section 3), obtaining c_{t+1} requires no separate pre-shifted copy of the action-probability arrays. The same reverse-order load, offset by one position, provides $\pi(a_{t+1} | s_{t+1})$ and $\mu(a_{t+1} | s_{t+1})$ directly in registers alongside $\pi(a_t | s_t)$ and $\mu(a_t | s_t)$ (Section 4.2).

The boundary carry is fixed at zero, $\Delta_T = 0$. The expected Q-values at s_T are already encoded in α_{T-1} . Moreover, $\beta_{T-1} = \gamma c_T(1 - d_{T-1})$ contains c_T , which would require the behaviour-policy probability of an action at s_T that was never sampled; we therefore take $c_T = 0$. Interior values Δ_t are nonzero in general and are computed by the scan.

3.3.4 TD(λ) Returns and Discounted Returns

TD(λ) returns [Sutton and Barto, 2018] interpolate between one-step TD ($\lambda = 0$) and Monte Carlo ($\lambda = 1$):

$$G_t^\lambda = r_t + \gamma[(1 - \lambda)V(s_{t+1}) + \lambda G_{t+1}^\lambda].$$

Expanding and grouping into $\alpha_t + \beta_t \cdot G_{t+1}^\lambda$ form:

$$\alpha_t = r_t + \gamma(1 - \lambda)(1 - d_t^{\text{term}})V(s_{t+1}), \quad \beta_t = \gamma\lambda(1 - d_t).$$

Discounted returns [Sutton and Barto, 2018] are the special case $\lambda = 1$ with no value function: $\alpha_t = r_t$, $\beta_t = \gamma(1 - d_t)$. Two boundary cases require separate treatment. At *interior truncated steps* ($d_t^{\text{trunc}} = 1$, $t < T-1$), $\beta_t = 0$ severs the carry, so the supplied continuation value must enter α_t directly. For TD(λ), this gives $\alpha_t = r_t + \gamma V(s_{t+1}^{\text{true}})$. At the *window boundary* ($t = T-1$, continuing episode), $\beta_{T-1} = \gamma\lambda$ and the bootstrap is passed as the scan carry $G_T = V(s_T)$, entering through both terms at once:

$$G_{T-1} = \underbrace{r_{T-1} + \gamma(1 - \lambda)V(s_T)}_{\alpha_{T-1}} + \gamma\lambda V(s_T) = r_{T-1} + \gamma V(s_T).$$

The two coefficients on $V(s_T)$ total γ , pricing the bootstrap in exactly once.

3.3.5 Eligibility Traces

Eligibility traces are one of the two **forward** recurrences in the library:

$$\mathbf{z}_t = \gamma\lambda(1 - d_{t-1})\mathbf{z}_{t-1} + \nabla_{\mathbf{w}}\hat{V}(s_t, \mathbf{w}_t), \quad d_{-1} := 0.$$

This maps to $\alpha_t = \nabla_{\mathbf{w}}\hat{V}(s_t, \mathbf{w}_t)$, $\beta_t = \gamma\lambda(1 - d_{t-1})$: the carry into t is severed by the *preceding* step’s flag, since d_t marks where an episode ends, not where the next one starts. Unlike the backward algorithms, the sequence is loaded in *chronological* order ($t = 0 \rightarrow T-1$, past to present) so each position accumulates a weighted sum of past gradients rather than future TD errors. The same combine function applies; only the load order differs. For linear function approximation, $\nabla_{\mathbf{w}}\hat{V}(s_t, \mathbf{w}_t) = \mathbf{x}_t$ (the feature vector).

3.3.6 Episodic Prefix Sum

The episodic prefix sum (segmented scan) is the second **forward** scan and computes a running total that resets at segment boundaries. Two boundary conventions are supported via a boundary argument, depending on whether the flag marks the end of the preceding segment or the start of the new one. The default, `boundary="ends_at"`, matches the rollout convention used by the other kernels in this library:

$$C_t = x_t + (1 - d_{t-1})C_{t-1}, \quad d_{-1} := 0,$$

mapping to $\alpha_t = x_t$, $\beta_t = 1 - d_{t-1}$. Passing `boundary="starts_at"` instead resets at t ($\beta_t = 1 - d_t$), for callers whose boundary flag marks a new segment’s first element directly rather than an ending one.

4 Implementation

4.1 Kernel Architecture

Each environment row is processed by one Triton program instance, indexed by `tl.program_id(0)`. Within the program, a block of `BLOCK_SIZE` timestep elements is distributed across the program’s warps and processed in registers and on-chip memory [Kirk and Hwu, 2016]. Thus, parallelism is exposed both across environment rows through the launch grid and across timesteps within each program.

The generic backward kernel consumes precomputed α and β tensors and is used by the chunked fallback in Section 4.3. For `seq_len ≤ 131072`, the algorithm-specific fused kernels instead load the raw rollout tensors and construct α_t and β_t in-register, avoiding their materialization in HBM. Both paths then follow the same scan procedure:

1. **Load** (HBM $\rightarrow$ registers): Load the required inputs for `BLOCK_SIZE` timesteps, in reverse order for backward scans, using `tl.arange`-based offsets. Masked loads handle positions beyond the actual sequence length.
2. **Construct** α, β (registers): Fused kernels derive α_t and β_t from the raw inputs; the generic kernel receives them directly.
3. **Scan** (on-chip): Apply `tl.associative_scan((α, β), axis=0, combine_fn=combine)`, giving $O(\log \text{BLOCK_SIZE})$ dependency depth. Out-of-range positions are excluded by the load/store masks.
4. **Bootstrap** (registers): Incorporate the algorithm-dependent boundary carry A_T as `scan_result + decay_product * bootstrap` (Section 3.2).
5. **Store** (registers $\rightarrow$ HBM): Write the results in original time order using the reversed offsets.

Table 2: HBM accesses for GAE (with bootstrap values, no truncations). r : rewards; v : values; d : terminated; bs : bootstrap values; δ : TD residuals; β : decay scalars; A : advantages.

Step	Reads	Writes
<i>Fused pipeline (single Triton kernel)</i>		
Load r, v (read twice at offsets rev and $rev+1$), d, bs ; compute α, β in-register; write A	5	1
Total		6
<i>Unfused pipeline (three GPU operations)</i>		
Compute $\delta = r + \gamma v_{t+1}(1-d) - v_t$ (δ is GAE's α_t ; v_{t+1} reads bs at the boundary column, a 5th tensor)	5	1
Compute $\beta = \gamma\lambda(1-d)$	1	1
Scan: $(\delta, \beta) \rightarrow A$	2	1
Total		11

4.2 The Fused Kernel Pattern

Fusion eliminates the HBM materialization of α and β between preprocessing and the scan. Table 2 illustrates the resulting traffic reduction for GAE with bootstrap values; the same pattern applies to the other fused kernels.

For GAE, fusion reduces the counted HBM accesses from 11 to 6 by eliminating the write and subsequent read of δ and β . Compile-time `constexpr` flags further specialize common cases:

- `HAS_TRUNCATIONS=False`: skips the truncateds array and the 2D `bootstrap_values[num_envs, seq_len]` tensor (used for per-step truncation bootstraps), saving two HBM reads. The window-boundary bootstrap is still supported via a separate 1D scalar argument `bootstrap_ptr[num_envs]`, controlled by `HAS_BOOTSTRAP`.
- `HAS_BOOTSTRAP=False`: substitutes 0.0 for the window-boundary bootstrap and avoids materializing an all-zero `bootstrap_ptr`. Relative to the previous path, this reduces full-call time by $\sim 20\text{--}27\%$ at small shapes on both GPUs evaluated here.²

4.3 Chunked Fallback for Long Sequences

The flat kernel is limited to `seq_len` $\leq$ 131072 (the maximum `BLOCK_SIZE` supported by Triton's `tl.associative_scan`). For longer sequences, backward algorithms (GAE, V-Trace, Retrace, $\text{TD}(\lambda)$, discounted returns) fall back to a chunked PyTorch path that processes the sequence in blocks and propagates the carry between blocks. The eligibility trace and prefix-sum forward scans do not have a chunked fallback: a left-to-right chunked forward scan is structurally feasible but was not implemented because sequences longer than 131072 timesteps are very uncommon in on-policy RL, making the added implementation complexity hard to justify.

The chunked path is unfused: it materializes α and β as full tensors before invoking the scan. At `seq_len` $>$ 131072, throughput is dominated by global memory bandwidth rather than kernel launch overhead, so the fusion benefit is smaller and the additional implementation complexity of a fused chunked kernel is not justified for typical RL workloads.

4.4 Numerical Precision

All kernels require `float32` inputs and raise a `TypeError` on other dtypes. The scan accumulates T additions over potentially thousands of timesteps; `bfloat16`'s 7 mantissa bits cause meaningful numerical drift at large T compared to `float32`'s 23 mantissa bits. In mixed-precision training pipelines (`torch.autocast`), the recommended pattern is to cast advantage inputs to `float32` before the scan and allow the policy forward pass to run in `bfloat16`.

²The isolated benchmark compares the previous `torch.zeros(num_envs)` path with the current `HAS_BOOTSTRAP=False` path, saving $\sim 8\text{--}10\ \mu\text{s}$ (`benchmarks/measure_bootstrap_skip.py`).

5 Correctness Analysis

5.1 Algebraic Verification

Because $\oplus$ is associative (Section 2.2), the T per-timestep tuples can be reduced in any grouping without changing the result: a $\log_2(T)$ -depth tree reduction yields the same accumulation as the sequential scan, in $\log_2(T)$ dependent steps rather than T . Concretely, for a 4-step window the reduction ends with one thread holding a single accumulated tuple spanning all four timesteps; the α component of that tuple works out to

$$A_0 = \alpha_0 + \beta_0(\alpha_1 + \beta_1(\alpha_2 + \beta_2\alpha_3)),$$

the exact sequential expansion of $A_t = \alpha_t + \beta_t A_{t+1}$ – the value at the first chronological timestep, depending on all four local recurrence terms. Every combine pairs a chronologically-earlier tuple with a chronologically-later one, per Equation 2, so Section 5.2’s boundary-severing argument (stated directly in these same $t/t+1, \dots$ terms) applies here unchanged. Appendix A gives the full thread-by-thread trace.

5.2 Episode Boundary Correctness

When $d_t = 1$ at a boundary step, $\beta_t = 0$, and the operator severs the carry. Combining step t ’s own tuple $(\alpha_t, 0)$ with $(\alpha_{t+1..}, \beta_{t+1..})$ – the tuple already accumulated for steps $t+1, \dots$ – gives

$$(\alpha_t, 0) \oplus (\alpha_{t+1..}, \beta_{t+1..}) = (\alpha_t + 0 \cdot \alpha_{t+1..}, 0 \cdot \beta_{t+1..}) = (\alpha_t, 0).$$

Everything accumulated for later timesteps is multiplied by zero, and the resulting β is zero, so no credit propagates backward past t . The scan result at t is therefore unaffected by any following episode across a done flag.

The tree reduction therefore reproduces the sequential recurrence, while $\beta_t = 0$ prevents propagation across episode boundaries. For the two forward scans, the same argument applies with the shifted boundary coefficient: when $d_{t-1} = 1$, $\beta_t = 0$ prevents carry from the preceding episode from entering timestep t .

6 Benchmarking Methodology

6.1 Measurement Protocol

All benchmarks use CUDA event timing with the following protocol:

1. **Warmup** (20 iterations, untimed): absorbs Triton JIT compilation, autotuning, and first-touch allocation.
2. **Timed iterations** (50 per trial, 5 trials): each iteration is individually timed with CUDA event synchronization before start and after stop.
3. **Robust estimator**: each trial reports its median over 50 iterations; the final result is the minimum of 5 trial medians.

Credit-assignment errors are often finite and plausible-looking rather than crashes – for example, seeding the window carry with the bootstrap for GAE would double-count $V(s_T)$ at every position (Section 3.2) – so they survive testing that checks only for NaN/Inf. Every kernel is therefore validated against an independent sequential reference implementation (`atol=rtol=1e-4`) before timing results are included in the evaluation.

Following established parallel benchmarking practice [Hoeffler and Belli, 2015], we report the minimum of five independently warmed trial medians to reduce sensitivity to occasional interference from frequency changes or OS scheduling. Such fixed-overhead events disproportionately affect short kernels and can otherwise produce non-monotonic measurements across problem sizes.

All GPU benchmarks (both the H100 and the RTX 2000 Ada in Section 7) were run on RunPod-hosted cloud instances, on `torch==2.4.1+cu124` with Triton 3.0.0.

6.2 Baseline

The baseline uses a linear-space scan: `parallel_suffix_scan/parallel_prefix_scan`, a $\log_2(T)$ -doubling associative scan with no `log/exp` anywhere,³ verified the same way. All seven baselines share the same $\log_2(T)$ -doubling engine

³A log-space formulation based on cumulative sums of $\log(\beta_t)$ was also considered but rejected because repeated termination resets underflow in `float32`. At the termination rate used in these benchmarks, two to three terminations are sufficient to push the running sum below the representable range, producing `inf/nan` outputs. We therefore use the linear-space scan.

Table 3: Production-regime speedup (`num_envs=4096`, `seq_len=128`), full-call wall time, v0.1.3 release measurements. *Loop* is the Triton kernel’s speedup over the naive uncompiled sequential-loop baseline. *Plain* and *With trunc.* are its speedup over the vectorized `torch.compile` baseline (Section 6.2). Eligibility traces and prefix sum have no with-truncations path.

Algorithm	Loop	H100 80GB HBM3		Loop	RTX 2000 Ada	
		Plain	With trunc.		Plain	With trunc.
GAE	317×	2.36×	1.7×	281×	2.57×	1.6×
V-Trace	96×	2.78×	2.7×	86×	3.46×	3.3×
Retrace(λ)	72×	1.91×	1.9×	16×	1.62×	1.6×
λ -returns	267×	2.85×	2.6×	251×	5.14×	4.4×
Discounted returns	167×	2.70×	2.6×	171×	5.70×	4.6×
Eligibility traces	191×	2.48×	–	207×	2.28×	–
Prefix sum	167×	2.39×	–	168×	2.25×	–

– `parallel_suffix_scan` for the backward algorithms, `parallel_prefix_scan` (its direct forward-recurrence mirror) for eligibility traces and prefix sum – so the comparison uses one consistent baseline strength across algorithms.

Baseline scope The results below are relative to this implementation; they do not establish optimality among PyTorch implementations. The doubling-scan baseline pays 6–12 kernel launches per call (one per doubling step, growing with `seq_len`), against 1–2 launches for the `rl-triton` Triton kernel it is compared against. This launch-count gap is architectural: `torch.compile/Inductor` does not fuse this elementwise-then-scan pattern into as few kernel launches as a hand-written Triton kernel can.

Every “ $N\times$ vs. `torch.compile`” ratio reported in Section 7 is measured against this specific baseline.

6.3 Performance Regression Guard

The library includes a CI performance test that checks a GPU-specific minimum speedup over the vectorized baseline at 128 environments and 1024 steps. For each GPU, the threshold is set to 90% of the minimum observed across three independent runs. Across the currently calibrated H100, H200, and RTX 2000 Ada GPUs, the lowest threshold is $1.61\times$ (Retrace on RTX 2000 Ada).

7 Results

All results use the v0.1.3 benchmark data,⁴ measured under the protocol in Section 6.1 against the baseline in Section 6.2. We report results on an NVIDIA H100 80GB HBM3 and an NVIDIA RTX 2000 Ada Generation. Unless stated otherwise, speedups are full-call ratios for the complete `compute*(tensors) -> tensors` invocation, including launch and wrapper overhead. Device-only ratios are higher (1.76 – $13.19\times$ at the headline configuration versus 1.6 – $5.70\times$ full-call), so we report full-call ratios throughout.

7.1 Massively-Parallel-Simulation Headline

Table 3 reports all seven algorithms at `num_envs=4096`, `seq_len=128`: the massively-parallel-simulation regime this library targets (e.g., Isaac Gym/Isaac Lab-style environment counts at short rollout horizons). It reports three tiers per algorithm – the naive loop, the vectorized-and-compiled scan, and the fused Triton kernel – plus the with-truncations path (eligibility traces and prefix sum have no truncation-aware variant). A sequential NumPy CPU loop is also measured but not tabulated; it is several times slower than the GPU loop.⁵

The 16 – $317\times$ speedup over the naive loop primarily reflects the reduction from an $O(T)$ sequential recurrence to an $O(\log T)$ associative scan (Section 2.2). The vectorized `torch.compile` baseline already captures much of this gain. Relative to that baseline, fusion provides a further 1.6 – $5.70\times$ speedup, ranging from $1.6\times$ for Retrace to $5.70\times$ for

⁴Full per-shape tables, measurement scripts, and implementation notes are available with the library at <https://github.com/simonsays1980/rl-triton>.

⁵E.g. for GAE at `num_envs=4096`: 180ms (H100) / 56ms (RTX 2000 Ada) for the NumPy loop vs. 21ms / 9ms for the GPU loop.




Figure 1: Launch/wrapper overhead and device time for GAE (plain path, H100) across the twelve shapes in Table 4. *Top*: launch/wrapper overhead as a fraction of full-call time for `rl-triton` and the `torch.compile` baseline, decreasing from ~ 88 – 95% at the smallest shapes to ~ 11 – 37% at the largest. *Bottom*: full-call time decomposed into launch/wrapper and device components. `rl-triton`’s launch/wrapper cost remains ~ 0.026 – 0.035 ms, while the baseline’s is 1.4 – $3.1\times$ larger.

discounted returns (Table 3). Discounted returns and λ -returns sit at the high end because they are cheapest per step – no importance-sampling ratio, fewer value-function reads than GAE, V-Trace, or Retrace – so the kernel does less per-step work while the shared baseline pays the same cost regardless of algorithm. Figure 1 shows why this regime favors fusion.

The smaller speedup on the truncation-aware path arises from a property of the Triton kernel, not the baseline. The `torch.compile` baseline compiles and runs the with-truncations graph unconditionally, so its own time is essentially unchanged whether truncations are all-zero or real. The Triton kernel, by contrast, uses a `HAS.TRUNCATIONS` compile-time fast path (Section 4.2) that skips the truncateds array and the 2D bootstrap-values tensor when there are no truncations. For GAE, V-Trace, λ -returns, and discounted returns, the plain path is consistently faster in absolute terms and therefore generally achieves a larger speedup over `torch.compile`. Retrace shows little systematic difference between the two paths.

Retrace(λ) gives the smallest speedup at this shape on both GPUs, reflecting its heavier per-timestep work relative to the other algorithms. Its with-truncations ratio is still smaller than its plain ratio, like every other algorithm ($1.91\times \rightarrow 1.9\times$ on H100, $1.62\times \rightarrow 1.6\times$ on RTX 2000 Ada), but by a much smaller gap: the savings from the `HAS.TRUNCATIONS` fast path constitute a smaller fraction of its total runtime. The same effect also narrows its speedup over the sequential loop, whose HBM round-trips account for a smaller fraction of the heavier computation.

The relative H100 and RTX speedups are shape-dependent: they vary by algorithm and can reverse between shapes for the same algorithm (Section 9.2). Neither GPU consistently yields a larger Triton-to-`torch.compile` speedup.

This massively-parallel-simulation shape (`num_envs=4096`, `seq_len=128`) is short by design, and it is not where the kernel’s advantage is largest. The $O(\log T)$ -depth argument (Section 2.2) predicts the opposite: longer rollouts should

Table 4: Full shape sweep, plain path, full-call speedup vs. the vectorized `torch.compile` baseline (v0.1.3 release measurements).

num_envs	seq_len	H100 80GB HBM3		RTX 2000 Ada	
		GAE	Retrace	GAE	Retrace
64	512	2.7×	2.2×	2.7×	2.0×
128	1024	3.2×	2.3×	3.1×	1.8×
256	1024	3.3×	2.3×	3.0×	1.4×
512	128	2.8×	2.0×	2.5×	1.8×
512	512	2.9×	2.2×	2.8×	1.4×
512	2048	3.0×	1.4×	5.7×	2.0×
512	4096	3.2×	0.6×	6.3×	0.8×
4096	128	2.8×	1.9×	2.5×	1.6×
4096	512	3.3×	1.6×	5.3×	2.4×
4096	2048	5.3×	1.8×	7.8×	2.9×
16384	128	2.8×	1.9×	4.8×	2.3×
16384	512	5.1×	1.9×	6.8×	2.7×

benefit *more*, since the vectorized baseline pays proportionally more $O(\log T)$ HBM round-trips as T grows while the fused kernel avoids intermediate HBM round-trips between scan stages. We therefore evaluate sequence lengths up to 4096; Section 7.2 reports the resulting scaling behavior.

7.2 Full-Grid Scaling: GAE and Retrace

Table 4 reports the full shape sweep (twelve $(\text{num_envs}, \text{seq_len})$ configurations, spanning $\text{num_envs} \in \{64, \dots, 16384\}$ and $\text{seq_len} \in \{128, \dots, 4096\}$) for GAE and $\text{Retrace}(\lambda)$ on both GPUs, plain path, full-call ratio. The corresponding sweeps for the remaining five algorithms are reported in Appendix B.

GAE remains above $2.5\times$ across the sweep and reaches its largest ratios at the larger shapes. Against the uncompiled sequential-loop baseline instead (not tabulated here at every shape; see Table 3 for one representative shape), GAE’s gain grows from $1464\times$ at $(64, 512)$ to $7570\times$ at $(512, 4096)$ on H100. The loop baseline executes an $O(T)$ serial chain of GPU operations, with repeated kernel launches and global-memory traffic at each timestep, whereas the associative scan has only $O(\log T)$ dependency depth. Retrace behaves differently: at $\text{seq_len}=4096$, its speedup over the doubling-scan baseline falls below $1\times$ ($0.6\times$ on H100 and $0.8\times$ on RTX 2000 Ada). It remains faster than the uncompiled sequential loop at every shape in the grid; its minimum loop-baseline speedup is $43.5\times$ on H100 and $4.6\times$ on RTX 2000 Ada. Section 7.3 analyzes the source of the regression.

7.3 Retrace’s Register-Pressure Regression Above $\text{seq_len}=2048$

$\text{Retrace}(\lambda)$ ’s fused kernel reads the $[\text{num_envs}, \text{seq_len}, \text{num_actions}]$ action-probability tensor a second time in-kernel to construct c_{t+1} , leaving both the original and shifted values live during the scan. At $\text{seq_len}=4096$, this increases register demand enough to cause spilling: spilling: at $\text{seq_len}=4096$ ($\text{num_envs}=512$, $\text{num_actions}=4$), `ptxas` reports 128 registers per thread, a 400-byte per-thread stack frame, and 456 bytes each of spill stores and spill loads, at 25% occupancy.⁶

The implementation therefore dispatches longer sequences to the generic materialize-then-scan path, which computes the 3D probability terms outside the scan kernel and avoids the same register-pressure limit. The $\text{seq_len}=2048$ threshold is conservative rather than the crossover between the two `rl-triton` paths: at $\text{seq_len}=4096$ the fused path is still faster than the generic path ($0.73\times$ vs. $0.47\times$ relative to `torch.compile`), whereas the generic path overtakes it at $\text{seq_len}=8192$.

⁶`ptxas` is NVIDIA’s PTX-to-SASS assembler and performs register allocation during compilation; registers that cannot be allocated are spilled to local memory, with spill traffic reported in bytes. On H100, each SM provides 64K 32-bit registers and supports at most 64 resident warps. At 128 registers per thread, one 32-thread warp consumes 4096 registers, limiting an SM to 16 resident warps, or 25% occupancy from register pressure alone. The spill traffic does not enter this occupancy calculation. Measurements use Triton 3.0.0 and `ptxas` 12.4.99 targeting `sm_90a`; register allocation is compiler-version dependent (`benchmarks/measure_retrace_register_pressure.py`).

Table 5: GAE share of a PPO update and end-to-end Triton speedup over each baseline, by policy hidden size (`seq_len=128`).

Hidden size	num_envs	GAE share / E2E vs. loop	GAE share / E2E vs. scan
(1024, 1024)	4096	$\sim 2.2\%$ / $\sim 1.02\times$	$< 0.1\%$ / $\sim 1.00\times$ (within noise)
(128, 128)	16384	$\sim 10\text{--}15\%$ / $\sim 1.11\text{--}1.16\times$	not measured at this configuration [†]

[†]The (128, 128) configuration was evaluated only against the loop baseline.

7.4 End-to-End PPO Speedup

We evaluate whether the kernel-level speedups translate into end-to-end PPO speedups using a synthetic update (`seq_len=128`, 4 epochs $\times$ 4 minibatches) on an H100. We compare against both the sequential backward loop used by CleanRL, RLlib, and Sample Factory (Section 8) and the vectorized `torch.compile` baseline of Section 6.2, at two policy sizes:

At the (1024, 1024) policy size, GAE accounts for $\sim 2.2\%$ of the PPO update against the loop baseline and less than 0.1% against the vectorized scan. The corresponding end-to-end speedups are $\sim 1.02\times$ and $\sim 1.00\times$, respectively, despite the larger isolated-kernel gains in Table 3. At (128, 128) ($\sim 150k$ parameters), GAE accounts for 10–15% of the update against the loop baseline, and the end-to-end speedup rises to $1.11\text{--}1.16\times$ across nine reruns.

These results closely follow Amdahl’s law: eliminating a component that accounts for 2.2% of an update can improve total runtime by at most $\sim 1.02\times$, whereas a 10–15% share permits approximately $1.11\text{--}1.18\times$. The microbenchmark results in Sections 7.1–7.2 should therefore be interpreted as isolated-kernel speedups rather than equivalent gains in training throughput.

For feed-forward policies, the addressable share can increase with longer rollouts because policy computation batches timesteps independently whereas credit assignment retains a temporal dependency; repeated advantage computation increases it further. Credit assignment is also learner-side work in high-throughput actor–learner methods such as IMPALA [Espoholt et al., 2018] and APPO [Petrenko et al., 2020], while PufferLib [Suarez, 2025b] similarly targets a small-policy, high-throughput regime with a custom CUDA advantage kernel [Suarez, 2025a]. At (128, 128), CUDA graph replay improves the optimizer step by only $\sim 1.02\text{--}1.08\times$, indicating that launch overhead is not the dominant cost.

7.5 Comparison Against PufferLib

We compare `rl-triton` against PufferLib’s hand-written CUDA advantage kernel (`puff_advantage_row_cuda`), which assigns one thread per environment row and performs a sequential $O(T)$ scan. Measurements were collected on an H100 after verifying both kernels against an independent reference on all jointly supported outputs.⁷

We evaluate two workload regimes. The *classic on-policy* regime (`seq_len` $\in \{128, \dots, 4096\}$, `num_envs` $\in \{128, \dots, 8192\}$, Figure 2) represents moderate-to-long rollouts with moderate parallelism, as commonly used in PPO-style training. The *massively parallel simulation* regime (`seq_len` $\in \{8, \dots, 128\}$, `num_envs` $\in \{4096, \dots, 32768\}$, Figure 3) covers short rollouts with high environment parallelism. This regime appears in systems such as PufferLib [Suarez, 2025b], Gigaflow [Cusumano-Towner et al., 2025], GPUdriVe [Kazemkhani et al., 2024], Nocturne [Vinitzky et al., 2022], Isaac Gym/Isaac Lab [Makoviychuk et al., 2021, NVIDIA Isaac Lab Team, 2025], and robot-locomotion PPO workloads surveyed by UniLab [Jia et al., 2026].

Across both regimes, `rl-triton` has lower full-call time at 31 of 40 evaluated shapes. Sequence length largely determines the crossover: in the massively parallel simulation regime, changing `num_envs` does not change the winner at any tested `seq_len` except $(\text{num_envs}, \text{seq_len}) = (32768, 32)$ (Figure 3). PufferLib exposes a single `done` flag per step and provides neither termination/truncation semantics nor a bootstrap-value mechanism, so these cases cannot be compared directly.

⁷The benchmark uses a pinned, sha256-verified copy of PufferLib 3.0.0’s CUDA extension source, JIT-compiled (`benchmarks/pufferlib_ext/`). Timing uses CUDA events with warmup, explicit synchronization, and the minimum of 11 per-trial medians (100 iterations/trial) per configuration; both kernels are checked against an independent reference before any timing is trusted. Full methodology, the equivalence proof, and the corresponding V-Trace comparison are provided with the `rl-triton` benchmark suite.




Figure 2: Full-call GAE performance of `rl-triton` and `PufferLib` in the classic on-policy regime on an H100. `rl-triton` has lower runtime across all 20 evaluated shapes.




Figure 3: Full-call GAE performance of `rl-triton` and `PufferLib` in the massively parallel simulation regime on an H100. `rl-triton` has lower runtime at 11 of 20 shapes, including every configuration with `seq_len` ≥ 32 except $(\text{num_envs}, \text{seq_len}) = (32768, 32)$; `PufferLib` is faster at all configurations with `seq_len` ≤ 16 . The crossover shifts to longer sequences as `num_envs` increases, reflecting `PufferLib`’s low overhead for very short per-thread sequential scans.

8 Related Work

RL frameworks. CleanRL [Huang et al., 2022], RLlib [Liang et al., 2018], and Sample Factory [Petrenko et al., 2020] implement GAE and return computation in their training stacks. Representative PyTorch implementations evaluate the temporal recurrence with sequential backward loops rather than an associative scan. `rl-triton` provides GPU implementations with distinct handling of termination, truncation, and rollout-window bootstrapping (Section 3.2).

Parallel scans in ML. Associative scans are widely used to parallelize linear recurrences in sequence models. State-space and linear recurrent architectures such as S4 [Gu et al., 2022], S5 [Smith et al., 2023], LRU [Orvieto et al., 2023], and Mamba [Gu and Dao, 2023] exploit related parallel recurrence formulations. LRUs have also been used as recurrent backbones for partially observable RL. `rl-triton` instead applies associative scans to post-rollout scalar credit assignment rather than hidden-state sequence modeling.

IO-aware GPU kernels. FlashAttention [Dao et al., 2022, Dao, 2023] demonstrates how tiling and fusion around the GPU memory hierarchy can reduce HBM traffic and accelerate memory-intensive ML operators. `rl-triton` applies the same IO-aware principle to RL credit assignment, using Triton to keep the associative scan on-chip.

Off-policy correction. V-Trace [Espeholt et al., 2018] and Retrace(λ) [Munos et al., 2016] are the off-policy correction algorithms implemented in `rl-triton`. ACER [Wang et al., 2017] uses related truncated importance-sampling corrections with trust-region stabilization but is not currently implemented. These methods contain recursive importance-weighted return computations to which the affine-scan formulation applies.

9 Discussion and Future Work

9.1 Applicability Beyond Standard RL

The recurrence implemented by `rl-triton` also appears outside conventional on-policy RL, although direct applications beyond RL are limited.

- **Long-context LLM post-training.** PPO-style LLM training applies credit assignment over generated tokens, so long-chain-of-thought responses can turn advantage estimation into a thousands-step recurrence. T-PPO computes GAE-style advantages over 8K-token truncated windows within responses of up to 24K tokens [Fan et al., 2025], while recent long-context PPO experiments evaluate token-level GAE with up to 8192-token responses and 4096 rollouts per update [Gong et al., 2026]. The same opportunity is visible in current LLM-RL implementations: SkyRL [Cao et al., 2025] and `verl` [Sheng et al., 2024] implement token-level GAE with reverse sequential loops, while their REINFORCE++ estimators likewise compute discounted returns by scanning backward over response tokens. These workloads map directly to the fused GAE and discounted-return scans implemented here.
- **VLA reinforcement learning.** PPO-based post-training of vision-language-action policies requires advantage estimation over environment interaction trajectories. RLinf-VLA [Zang et al., 2025], for example, supports PPO training of VLA policies across robotic simulators, and its current implementation computes GAE over $[T, B]$ rollout tensors with a reverse Python loop over T . This recurrence maps directly to the fused GAE scan implemented here, making VLA training with many parallel environments and longer interaction horizons another natural application.
- **World-model reinforcement learning.** DreamerV4 [Hafner et al., 2025] trains its value function on imagined trajectories using bootstrapped λ -returns and derives its policy advantages from the same targets. These returns obey exactly the backward affine recurrence implemented here, making imagination-based world-model training another direct application of the fused λ -return scan.
- **Spiking-neural-network BPTT.** A direct non-RL instance occurs in backpropagation through hard-reset spiking neurons when reset differentiation is detached. The forward pass generates a runtime spike sequence S_t from the neuron state; spike times, and hence the resulting segment lengths, are not known in advance. The backward temporal gradient then has the affine form

$$g_t = \alpha_t + \beta_t g_{t+1}, \quad \beta_t = \kappa_t (1 - S_t),$$

so each spike zeros the temporal carry. SNN frameworks such as SpikingJelly provide dedicated GPU kernels for multi-step neuron forward and backward passes [Fang et al., 2023], making this a GPU-relevant instance of the same segmented affine recurrence.

Direct applicability requires a scalar affine recurrence and enough independent sequences for GPU parallelism. RL credit assignment additionally attaches reset and bootstrap semantics to per-step boundary signals. Hard-reset spiking-neural-network BPTT with detached reset differentiation is a non-RL example with comparable recurrence-level boundary semantics.

9.2 Limitations

Retrace register pressure above `seq_len=2048`. `Retrace( $\lambda$ )` becomes slower than `torch.compile` above `seq_len=2048` because of the register-pressure effects described in Section 7.3.

GAE device-time regression at the shortest, highest-parallelism shape. At `num_envs=16384` and `seq_len=16` on H100, GAE’s device-only ratio falls below $1\times$ relative to `torch.compile`, while its full-call ratio remains above $1\times$ because launch and wrapper overhead dominate at this shape. GAE is the only one of the five measured algorithms that exhibits this inversion.

Warp-count tuning gap above `BLOCK_SIZE=16384`. Per-`BLOCK_SIZE` warp counts are tuned only up to `BLOCK_SIZE=16384`; larger, reachable sizes fall back to an untuned default of 16 warps. Tuning at small `BLOCK_SIZE` recovered $2.1\text{--}2.7\times$ over an untuned default there, suggesting that larger block sizes may also benefit from tuning. We leave this evaluation to future work.

The H100-vs-RTX margin direction is shape-dependent – open question. At a small, launch-overhead-dominated shape (128 environments, 1024 steps) vs. the massively-parallel-simulation shape of Table 3, GAE, V-Trace, and λ -returns each reverse which card shows the larger Triton-vs-torch.compile margin; Retrace alone favors H100 at both. This reversal is a shape effect rather than evidence of a general "consumer vs. datacenter" trend; its underlying mechanism remains unidentified.

Baseline dependence. All reported speedups are relative to the baseline described in Section 6.2. A faster numerically stable PyTorch implementation would reduce these ratios.

Sequence length limit. Every flat fused kernel is limited to `seq_len` $\leq$ 131072 by Triton’s internal block size constraints; the backward kernels fall back to an unfused, two-pass chunked path beyond that, but the two forward scans (eligibility traces, episodic prefix sum) have no chunked fallback at all.

Float32 only. All kernels require `float32` inputs (Section 4.4); `bfloat16` support is deferred to a future release pending a precision-tradeoff analysis across algorithms and sequence lengths.

Scalar β only. β_t must be a scalar per timestep, shared across whatever is scanned. Per-dimension β_t , as true online TD(λ)/per-parameter eligibility traces require [van Seijen et al., 2016, Sutton and Barto, 2018], is not supported: β_t is baked into the kernel launch as a scalar rather than read per-lane from memory, and supporting it would require a different kernel structure.

9.3 Planned Extensions

Planned work focuses on broader rollout support and lower-overhead execution. For multi-turn and tool-use trajectories, we plan observation-skip masking that carries the recurrence across masked tokens ($\beta_t = 1$), together with fused advantage normalization and value loss. Systems work includes `bfloat16` I/O, chunked forward scans, CUDA graph support, and fixing the Retrace register-pressure regression (Section 7.3). We also plan differentiable PPO, GRPO, and KL-loss kernels with hand-written backward passes.

10 Conclusion

The seven credit-assignment algorithms considered here – GAE, V-Trace, Retrace(λ), TD(λ) returns, discounted returns, eligibility traces, and episodic prefix sums – share a common first-order linear recurrence structure that can be evaluated with the same associative scan operator. Each algorithm is implemented by its own fused Triton kernel, which constructs its recurrence coefficients on-chip. By keeping the $O(\log T)$ scan on-chip, `rl-triton` avoids the sequential loop’s per-timestep HBM round-trips and the vectorized `torch.compile` baseline’s intermediate HBM round-trips between doubling stages. In the massively parallel simulation regime, this yields 1.6–5.70 $\times$ full-call speedups over the vectorized baseline. For most algorithms, the gains increase at longer sequence lengths; Retrace is the exception because of its long-sequence register-pressure regression. The implementation preserves the specified handling of terminated episodes, truncated episodes, and rollout-window boundaries.

Acknowledgements

The author thanks Artur Niederfahrenheit (Anyscale) for valuable feedback and discussions, and Mark Towers (Gymnasium maintainer) for comments on an early draft and support with the arXiv submission.

References

- Guy E. Blelloch. Prefix sums and their applications. Technical Report CMU-CS-90-190, School of Computer Science, Carnegie Mellon University, 1990. 2
- Shiyi Cao, Sumanth Hegde, Dacheng Li, Tyler Griggs, Shu Liu, Eric Tang, Jiayi Pan, Xingyao Wang, Akshay Malik, Graham Neubig, Kourosh Hakhmaneshi, Richard Liaw, Philipp Moritz, Matei Zaharia, Joseph E. Gonzalez, and Ion Stoica. SkyRL-v0: Train real-world long-horizon agents via reinforcement learning, 2025. 14
- Marco Cusumano-Towner, David Hafner, Alex Hertzberg, Brody Huval, Aleksei Petrenko, Eugene Vinitsky, Erik Wijmans, Taylor Killian, Stuart Bowers, Ozan Sener, Philipp Krähenbühl, and Vladlen Koltun. Robust autonomy emerges from self-play. *arXiv preprint*, 2025. arXiv:2502.03349. 12

- Tri Dao. FlashAttention-2: Faster attention with better parallelism and work partitioning. *arXiv preprint arXiv:2307.08691*, 2023. 13
- Tri Dao, Daniel Y. Fu, Stefano Ermon, Atri Rudra, and Christopher Ré. FlashAttention: Fast and memory-efficient exact attention with IO-awareness. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. arXiv:2205.14135. 3, 13
- Lasse Espeholt, Hubert Soyer, Rémi Munos, Karen Simonyan, Volodymyr Mnih, Tom Ward, Yotam Doron, Vlad Firoiu, Tim Harley, Iain Dunning, Shane Legg, and Koray Kavukcuoglu. IMPALA: Scalable distributed deep-RL with importance weighted actor-learner architectures. In *International Conference on Machine Learning (ICML)*, 2018. arXiv:1802.01561. 1, 5, 12, 13
- Tiantian Fan, Lingjun Liu, Yu Yue, Jiaze Chen, Chengyi Wang, Qiying Yu, Chi Zhang, Zhiqi Lin, Ruofei Zhu, Yufeng Yuan, Xiaochen Zuo, Bole Ma, Mofan Zhang, Gaozhong Liu, Ru Zhang, Haotian Zhou, Cong Xie, Ruidong Zhu, Zhi Zhang, Xin Liu, Mingxuan Wang, Lin Yan, and Yonghui Wu. Truncated proximal policy optimization. *arXiv preprint arXiv:2506.15050*, 2025. 14
- Wei Fang, Yanqi Chen, Jianhao Ding, Zhaofei Yu, Timothée Masquelier, Ding Chen, Liwei Huang, Huihui Zhou, Guoqi Li, and Yonghong Tian. SpikingJelly: An open-source machine learning infrastructure platform for spike-based intelligence. *Science Advances*, 9(40):ead1480, 2023. 14
- Xue Gong, Qi Yi, Ziyuan Nan, Guanhua Huang, Kejiao Li, Yuhao Jiang, Ruibin Xiong, Zenan Xu, Jiaming Guo, Shaohui Peng, and Bo Zhou. Segmental advantage estimation: Enhancing PPO for long-context LLM training. *arXiv preprint arXiv:2601.07320*, 2026. 14
- Albert Gu and Tri Dao. Mamba: Linear-time sequence modeling with selective state spaces. *arXiv preprint arXiv:2312.00752*, 2023. 13
- Albert Gu, Karan Goel, and Christopher Ré. Efficiently modeling long sequences with structured state spaces. In *International Conference on Learning Representations (ICLR)*, 2022. arXiv:2111.00396. 13
- Danijar Hafner, Wilson Yan, and Timothy Lillicrap. Training agents inside of scalable world models. *arXiv preprint arXiv:2509.24527*, 2025. 14
- Torsten Hoefler and Roberto Belli. Scientific benchmarking of parallel computing systems: Twelve ways to tell the masses when reporting performance results. In *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis (SC’15)*, 2015. 8
- Shengyi Huang, Rousslan Fernand Julien Dossa, Chang Ye, Jeff Braga, Dipam Chakraborty, Kinal Mehta, and João G. M. Araújo. CleanRL: High-quality single-file implementations of deep reinforcement learning algorithms. *Journal of Machine Learning Research*, 23(274):1–18, 2022. 13
- Yufei Jia, Zhanxiang Cao, Mingrui Yu, Heng Zhang, Shenyu Chen, Dixuan Jiang, Meng Li, Xiaofan Li, Yiyang Liu, Junzhe Wu, Zheng Li, XiLin Fang, Ting-Yu Tsui, Shengcheng Fu, Haoyang Li, Anqi Wang, Zifan Wang, Dongjie Zhu, Chenyu Cao, Zhenbiao Huang, Ziang Zheng, Jie Lu, Xin Ma, Zhengyang Wei, Xiang Zhao, Tianyue Zhan, Ye He, Yuxiang Chen, Yizhou Jiang, Yue Li, Haizhou Ge, Yuhang Dong, Fan Jia, Ziheng Zhang, Meng Zhang, Xiwa Deng, Zhixing Chen, Hanyang Shao, Chenxin Dong, Yixuan Li, Yizhi Chen, Bokui Chen, Kaifeng Zhang, Hanqing Cui, Yusen Qin, Ruqi Huang, Lei Han, Tiancai Wang, Xiang Li, Yue Gao, and Guyue Zhou. UniLab: A heterogeneous architecture for robot RL beyond GPU-dominant paradigms. *arXiv preprint*, 2026. arXiv:2605.30313. 12
- Saman Kazemkhani, Aarav Pandya, Daphne Cornelisse, Brennan Shacklett, and Eugene Vinitzky. GPUDrive: Data-driven, multi-agent driving simulation at 1 million FPS. *arXiv preprint*, 2024. arXiv:2408.01584. 12
- David B. Kirk and Wen-mei W. Hwu. *Programming Massively Parallel Processors: A Hands-On Approach*. Morgan Kaufmann, 3rd edition, 2016. 6
- Richard E. Ladner and Michael J. Fischer. Parallel prefix computation. *Journal of the ACM*, 27(4):831–838, 1980. 2
- Eric Liang, Richard Liaw, Robert Nishihara, Philipp Moritz, Roy Fox, Ken Goldberg, Joseph Gonzalez, Michael Jordan, and Ion Stoica. RLlib: Abstractions for distributed reinforcement learning. In *International Conference on Machine Learning (ICML)*, 2018. arXiv:1712.09381. 13
- Viktor Makoviychuk, Lukasz Wawrzyniak, Yunrong Guo, Michelle Lu, Kier Storey, Miles Macklin, David Hoeller, Nikita Rudin, Arthur Allshire, Ankur Handa, and Gavriel State. Isaac gym: High performance GPU-based physics simulation for robot learning. *arXiv preprint*, 2021. arXiv:2108.10470. 12
- Rémi Munos, Tom Stepleton, Anna Harutyunyan, and Marc G. Bellemare. Safe and efficient off-policy reinforcement learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2016. arXiv:1606.02647. 1, 5, 13

- NVIDIA Isaac Lab Team. Isaac lab: A GPU-accelerated simulation framework for multi-modal robot learning. <https://arxiv.org/abs/2511.04831>, 2025. arXiv:2511.04831. 12
- Antonio Orvieto, Samuel L. Smith, Albert Gu, Anushan Fernando, Caglar Gulcehre, Razvan Pascanu, and Soham De. Resurrecting recurrent neural networks for long sequences. In *International Conference on Machine Learning (ICML)*, 2023. arXiv:2303.06349. 13
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul Christiano, Jan Leike, and Ryan Lowe. Training language models to follow instructions with human feedback. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. arXiv:2203.02155. 1
- Aleksei Petrenko, Zhehui Huang, Tushar Kumar, Gaurav Sukhatme, and Vladlen Koltun. Sample factory: Egocentric 3D control from pixels at 100000 FPS with asynchronous reinforcement learning. In *International Conference on Machine Learning (ICML)*, 2020. arXiv:2006.11751. 1, 12, 13
- John Schulman, Philipp Moritz, Sergey Levine, Michael Jordan, and Pieter Abbeel. High-dimensional continuous control using generalized advantage estimation. In *International Conference on Learning Representations (ICLR)*, 2016. arXiv:1506.02438. 1, 2, 5
- John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017. 1
- Guangming Sheng, Chi Zhang, Zilingfeng Ye, Xibin Wu, Wang Zhang, Ru Zhang, Yanghua Peng, Haibin Lin, and Chuan Wu. HybridFlow: A flexible and efficient RLHF framework. *arXiv preprint arXiv:2409.19256*, 2024. 14
- Jimmy T. H. Smith, Andrew Warrington, and Scott W. Linderman. Simplified state space layers for sequence modeling. In *International Conference on Learning Representations (ICLR)*, 2023. arXiv:2208.04933. 13
- Joseph Suarez. PufferLib 3.0 release notes. <https://puffer.ai/blog.html>, 2025a. Accessed 2026-08-11. 12
- Joseph Suarez. PufferLib 2.0: Reinforcement learning at 1m steps/s. *Reinforcement Learning Journal*, 6:1378–1388, 2025b. 12
- Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. MIT Press, 2nd edition, 2018. 5, 6, 15
- Philippe Tillet, Hsiang-Tsung Kung, and David Cox. Triton: An intermediate language and compiler for tiled neural network computations. In *Proceedings of the 3rd ACM SIGPLAN International Workshop on Machine Learning and Programming Languages (MAPL)*, 2019. 3
- Harm van Seijen, A. Rupam Mahmood, Pilarski Christos, Oussama Mourad, and Richard S. Sutton. True online temporal-difference learning. *Journal of Machine Learning Research*, 17(145):1–40, 2016. 15
- Eugene Vinitsky, Nathan Lichtlé, Xiaomeng Yang, Brandon Amos, and Jakob Foerster. Nocturne: A scalable driving benchmark for bringing multi-agent learning one step closer to the real world. *arXiv preprint*, 2022. arXiv:2206.09889. 12
- Ziyu Wang, Victor Bapst, Nicolas Heess, Volodymyr Mnih, Rémi Munos, Koray Kavukcuoglu, and Nando de Freitas. Sample efficient actor-critic with experience replay. In *International Conference on Learning Representations (ICLR)*, 2017. arXiv:1611.01224. 13
- Samuel Williams, Andrew Waterman, and David Patterson. Roofline: An insightful visual performance model for multicore architectures. In *Communications of the ACM*, volume 52, pages 65–76, 2009. 3
- Hongzhi Zang, Mingjie Wei, Si Xu, Yongji Wu, Zhen Guo, Yuanqing Wang, Hao Lin, Peihong Wang, Liangzhi Shi, Yuqing Xie, Zhexuan Xu, Zhihao Liu, Kang Chen, Wenhao Tang, Quanlu Zhang, Weinan Zhang, Chao Yu, and Yu Wang. RLinf-VLA: A unified and efficient framework for reinforcement learning of vision-language-action models. *arXiv preprint arXiv:2510.06710*, 2025. 14

A Four-Thread Scan Trace

This appendix gives the complete thread-by-thread combine underlying the sketch in Section 5.1. The array is loaded in reverse chronological order, so thread i ($i = 1, \dots, 4$) initially holds the tuple $(\alpha_{4-i}, \beta_{4-i})$. We write $T_{a..b}$ for the accumulated tuple covering reversed positions a through b , and T_k as shorthand for the single-position tuple $T_{k..k}$ – thread k ’s value *before any reduction pass runs* (position 1 = last timestep $t=3$, position 4 = first timestep $t=0$):

$$T_1 = (\alpha_3, \beta_3), \quad T_2 = (\alpha_2, \beta_2), \quad T_3 = (\alpha_1, \beta_1), \quad T_4 = (\alpha_0, \beta_0).$$

Table 6: Full twelve-shape sweep for V-Trace, TD(λ)/ λ -returns, discounted returns, eligibility traces, and episodic prefix sum. Values are full-call speedups over the vectorized `torch.compile` baseline using the v0.1.3 release measurements.

H100 80GB HBM3							RTX 2000 Ada				
num_envs	seq_len	V-Trace	λ -ret.	Disc. ret.	Elig.	Prefix	V-Trace	λ -ret.	Disc. ret.	Elig.	Prefix
64	512	3.0×	3.0×	3.0×	2.8×	2.8×	2.9×	2.9×	2.9×	2.6×	2.5×
128	1024	3.0×	3.3×	3.3×	2.7×	2.7×	2.9×	3.1×	3.1×	2.6×	2.5×
256	1024	3.0×	3.2×	3.2×	2.8×	2.8×	2.9×	3.6×	3.3×	2.6×	2.6×
512	128	2.9×	2.9×	2.8×	2.5×	2.4×	2.8×	2.7×	2.6×	2.3×	2.2×
512	512	3.1×	3.1×	3.1×	2.8×	2.7×	3.0×	3.3×	3.3×	2.6×	2.5×
512	2048	3.0×	3.2×	3.3×	3.1×	3.2×	4.2×	16.8×	17.2×	4.9×	4.6×
512	4096	2.9×	5.9×	5.8×	3.2×	2.9×	6.2×	15.0×	24.5×	15.6×	15.1×
4096	128	2.8×	2.9×	2.8×	2.5×	2.5×	3.5×	5.5×	5.7×	2.3×	2.3×
4096	512	3.1×	4.9×	5.3×	2.7×	2.7×	5.7×	11.8×	27.9×	10.4×	8.8×
4096	2048	4.8×	8.5×	8.0×	5.4×	5.4×	7.2×	14.9×	18.1×	7.8×	7.9×
16384	128	3.0×	4.4×	4.2×	2.0×	2.2×	5.4×	10.0×	21.8×	8.5×	8.0×
16384	512	5.2×	9.0×	8.8×	4.8×	4.7×	6.6×	12.9×	15.6×	6.5×	6.5×

Reduction pass 1 (stride 1). Each thread reads from the thread exactly one position to its left. Thread 1 has no left neighbour and keeps its initial tuple. Threads 2, 3, and 4 each combine with their immediate left neighbour’s *initial* value T_k (all reads happen simultaneously before any writes, so every combine below uses the single-position tuples just defined, never an already-updated span):

$$\begin{aligned}
 T_{1..1} &= T_1 = (\alpha_3, \beta_3), \\
 T_{1..2} &= T_2 \oplus T_1 = (\alpha_2 + \beta_2\alpha_3, \beta_3\beta_2), \\
 T_{2..3} &= T_3 \oplus T_2 = (\alpha_1 + \beta_1\alpha_2, \beta_2\beta_1), \\
 T_{3..4} &= T_4 \oplus T_3 = (\alpha_0 + \beta_0\alpha_1, \beta_1\beta_0).
 \end{aligned}$$

After pass 1, threads 2, 3, and 4 each cover two adjacent positions.

Reduction pass 2 (stride 2). Each thread reads from the thread two positions to its left. Thread 4 combines its current tuple $T_{3..4}$ with thread 2’s result $T_{1..2}$:

$$T_{1..4} = T_{3..4} \oplus T_{1..2} = (\alpha_0 + \beta_0\alpha_1 + \beta_0\beta_1\alpha_2 + \beta_0\beta_1\beta_2\alpha_3, \beta_3\beta_2\beta_1\beta_0).$$

Thread 4 now covers all four positions, and the loop terminates – exactly $\log_2(\text{BLOCK_SIZE}) = \log_2 4 = 2$ fixed passes, no dynamic stopping criterion. Each thread now holds the correct tuple for its own position: thread 3’s, for instance, is $T_{2..3} \oplus T_{1..1} = T_{1..3}$ by the same stride-2 combine.

B Full-Grid Scaling: Remaining Algorithms

Table 4 (Section 7.2) reports the twelve-shape sweep for GAE and Retrace(λ), including Retrace’s long-sequence regression analyzed in Section 7.3. Table 6 reports the corresponding results for the remaining five algorithms.

Unlike Retrace, none of the remaining five algorithms falls below $1\times$ at any measured shape. At several long-sequence configurations, the RTX 2000 Ada ratios are substantially larger than on H100. This is driven mainly by the `torch.compile` baseline: for discounted returns at (512, 4096), the speedup is $5.8\times$ on H100 and $24.5\times$ on RTX, while the baseline is $11\times$ slower on RTX (2.62 ms vs. 0.24 ms) and the Triton kernel itself only $\sim 2.6\times$ slower.